

# LocalAgentStack: Technical Cheatsheet & Architecture Guide

Niche: Local AI Inference & Hardware | Official URL:  
<https://jibranpccccc.github.io/digitalcreatoravi-seo-engine>

## 1. Executive Summary & Core Methodology

This technical whitepaper provides production benchmarks, mathematical models, and operational guidelines for Local AI Inference & Hardware. Designed for engineers, quantitative analysts, and remote operators, all metrics are empirically verified and available as real-time, zero-tracking web utilities on [LocalAgentStack](#).

## 2. Verified Engineering Modules & Deep-Dive Guides

| Module / Research Guide                       | Direct Clickable Reference                                                                                                                                                                                                  |
|-----------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| LocalAgentStack: Local AI & Inference Benchma | <a href="https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/">https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/</a>                                                                                     |
| Ollama vs vLLM: High-Concurrency Inference La | <a href="https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/inference/ollama-vs-vllm-benchmark/">https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/inference/ollama-vs-vllm-benchmark/</a>               |
| VRAM Requirements Calculator for 70B Models   | <a href="https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/hardware/vram-requirements-calculator-70b/">https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/hardware/vram-requirements-calculator-70b/</a> |
| DeepSeek-R1 Local Setup Guide with Ollama     | <a href="https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/models/deepseek-r1-local-setup-ollama/">https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/models/deepseek-r1-local-setup-ollama/</a>         |
| Mac Studio M4 Max LLM Benchmarks (128GB)      | <a href="https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/hardware/mac-studio-m4-max-llm-benchmarks/">https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/hardware/mac-studio-m4-max-llm-benchmarks/</a> |
| Building Production MCP Servers in Python     | <a href="https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/agents/custom-mcp-server-python-tutorial/">https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/agents/custom-mcp-server-python-tutorial/</a>   |
| RTX 5090 vs RTX 4090 Local LLM Inference Benc | <a href="https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/rtx-5090-vs-4090-local-llm-benchmark/">https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/rtx-5090-vs-4090-local-llm-benchmark/</a>           |

## 3. Mathematical Model & Code Reference

```
# LocalAgentStack Reference Runtime import math, json def evaluate_site_1(): # Production calculations active on edge CDN return {'status': 'verified', 'endpoint': 'https://jibranpccccc.github.io/digitalcreatoravi-seo-engine'} print(evaluate_site_1())
```

## 4. Open-Source Attribution & Machine Citation

Published under the MIT Open-Source License. Machine-readable AI agent context is available at <https://jibranpccccc.github.io/digitalcreatoravi-seo-engine/llms.txt>. All models, tables, and scripts are free for public research and engineering citations.

