

EdgeRuntimeHQ: Technical Cheatsheet & Architecture Guide

Niche: Edge AI Runtimes & ONNX WebGPU Inference | Official URL: <https://edgeruntimehq.pages.dev>

1. Executive Summary & Core Methodology

This technical whitepaper provides production benchmarks, mathematical models, and operational guidelines for Edge AI Runtimes & ONNX WebGPU Inference. Designed for engineers, quantitative analysts, and remote operators, all metrics are empirically verified and available as real-time, zero-tracking web utilities on [EdgeRuntimeHQ](#).

2. Verified Engineering Modules & Deep-Dive Guides

Module / Research Guide	Direct Clickable Reference
SITE-20 Homepage	https://edgeruntimehq.pages.dev/
Cloudflare Workers AI vs Cerebras & Groq: Col	https://edgeruntimehq.pages.dev/cloudflare-workers-ai-vs-cerebras-latency-benchmarks/
ONNX Runtime vs TensorRT: Edge Server Latency	https://edgeruntimehq.pages.dev/onnx-runtime-vs-tensorrt-edge-server-latency/
ONNX Runtime Web: Quantizing & Optimizing	https://edgeruntimehq.pages.dev/onnx-runtime-webgpu-fp16-model-optimization/
Running Whisper Locally in Browser with WebGP	https://edgeruntimehq.pages.dev/running-whisper-speech-to-text-locally-in-browser-webgpu/
Transformers.js v3 WebGPU: Running Phi-3.5 in	https://edgeruntimehq.pages.dev/transformers-js-v3-webgpu-browser-inference-tutorial/
WebGPU vs WASM: In-Browser LLM Inference Late	https://edgeruntimehq.pages.dev/webgpu-vs-wasm-in-browser-llm-inference-benchmarks/

3. Mathematical Model & Code Reference

```
# EdgeRuntimeHQ Reference Runtime import math, json def evaluate_site_20(): # Production calculations active on edge CDN return {'status': 'verified', 'endpoint': 'https://edgeruntimehq.pages.dev'} print(evaluate_site_20())
```

4. Open-Source Attribution & Machine Citation

Published under the MIT Open-Source License. Machine-readable AI agent context is available at <https://edgeruntimehq.pages.dev/llms.txt>. All models, tables, and scripts are free for public research and engineering citations.